

Structura — Landing Page Content

CPL / ExtractionPlatform

July 27, 2026

CONTENTS

Structura	
Turn your most painful documents into clean, structured data.	
Hero	
The problem	
What we do	
How it works	
Any document. Any format. Any size.	
What makes us different	
Where this works	
How we work with you	
Why now	
Call to action	
FAQ	

Structura

TURN YOUR MOST PAINFUL DOCUMENTS INTO CLEAN, STRUCTURED DATA.

Working name — “Structura” is a placeholder. Swap for your final brand before launch.

Hero

Your critical data is trapped inside documents nobody can query.

Invoices. Benefit summaries. Regulatory filings. Spec sheets. Contracts. Rate cards. Policy manuals. The information your business runs on lives in thousands of pages of PDFs and scanned files that no dashboard can read and no analyst has time to key in by hand.

Structura turns those documents into structured, verified, queryable data. You hand us the messy source. We hand you clean, organized information you can trust, build on, and sell.

[[Book a pilot](#)] [[See how it works](#)]

The problem

Every operationally heavy industry has the same buried asset: the truth lives in documents, not databases.

- The document is dense, inconsistent, and written for humans, not machines.
- The volume is too high to key in by hand and too irregular for a simple template.
- The stakes are high enough that a wrong number is worse than no number.
- So the data stays locked up, and every team downstream rebuilds it from scratch, badly.

Companies pay analysts to read the same filings over and over. Product teams ship with stale numbers. Nobody has a single source of truth, because the source was never structured in the first place.

Off-the-shelf OCR gives you text, not data. Generic AI gives you a plausible guess with no way to tell right from wrong. Neither one is something you can put in front of a customer.

What we do

We build the structured layer your documents never had.

You give us the source material for a domain: a stack of PDFs, a set of filings, a batch of scanned forms, a list of web pages. We define the exact shape of the data that matters in your world, extract every field from every document, verify it, and deliver it in the form you actually work in: a clean spreadsheet, a database table, or a live feed, with a schema you own.

The result is not a pile of text. It is organized information where every entry is a real thing, every detail is clearly defined, and every value has been checked. It is the difference between “we have the documents” and “we have the data.”

How it works

- 1. We map your domain.** We sit with you and define the fields that matter: what an entity is, what each attribute means, what a valid value looks like. This becomes your schema. Full field names, precise types, no guessing about what a column holds.
- 2. We ingest the sources.** PDFs, web pages, scanned forms, photographed pages. Formats that fight back are our normal case, including sources that are geo-blocked, inconsistent between documents, or hundreds of megabytes each.
- 3. We extract into the schema.** Our engine reads each document and populates every field. It handles the fan-out (one plan across many counties, one contract across many line items) so the output is at the grain you actually query.

4. We verify and report the gaps. Every extraction is checked. You get a clean dataset plus a gap report: what was found, what the source genuinely never contained, and what needs review. You know exactly how much to trust every column before you build on it.

Re-runs are non-destructive and versioned. When a source updates, we re-extract without ever clobbering the data you already trust.

Any document. Any format. Any size.

The messiness of your source is our problem, not your line item. A pristine digital file and a stack of crooked scans come back as the same clean, organized result.

- **Text-based PDFs.** Born-digital files with selectable text. The easy case, done at volume.
- **Scanned and image-only PDFs.** Faxes, photocopies, signed forms, photographed pages. Documents that are pictures of text, not text. We read them anyway.
- **One page or ten thousand.** A single invoice or a thousand-page policy manual runs through the same pipeline.
- **Small, large, and enormous files.** Catalogs and filings that run to hundreds of megabytes and choke ordinary tools are a normal input, not a special request.
- **Mixed batches.** A folder where every document is a different layout, vintage, and quality. We normalize the whole pile into one consistent schema.
- **The parts that break templates.** Multi-column layouts, dense tables, stamps, handwriting, footnotes, forms that changed format three times over the years.

Whether a document is clean text or a wall of scanned faxes, you get the same structured, verified result. The extra difficulty is absorbed into the work. It never shows up as a caveat in your data or a gap in what you receive.

What makes us different

A schema you own, not a black box. The data comes out in a shape you defined, with real field names and types. It plugs straight into your database, your CRM, your ERP, your data warehouse, your custom product, or your customer's systems, without a cleanup project in between.

Verification is the product, not an afterthought. Anyone can get a language model to guess. We tell you which values were confirmed, which were absent from the source, and which need a human. That is what makes the data safe to sell.

Built generic from day one. The engine that breaks dense cost and coverage terms into clean, typed fields can be pointed at leases, filings, spec sheets, or invoices just as easily. New vertical, new schema, same machine.

We handle the ugly reality. Blocked sources, inconsistent formats, enormous files, entities that fan out across regions. The parts that make everyone else quote "custom project, call us" are our default path.

Cost and completeness under control. The engine is bounded so a hard document can't run up a runaway bill, and it sweeps the full field set so nothing quietly gets skipped.

Where this works

If an industry runs on documents, it needs Structura. A few of the North American verticals with exactly this shape:

- **Healthcare and benefits** — plan documents, formularies, provider directories, prior-auth rules.
- **Insurance (P&C, life, specialty)** — policy forms, endorsements, rate filings, coverage comparisons.
- **Financial services** — prospectuses, credit agreements, fund fact sheets, disclosure filings.
- **Finance and back office** — invoices, purchase orders, receipts, remittances, bank and card statements, expense reports across any vendor format.

- **Real estate** — leases, offering memoranda, appraisal reports, HOA and title documents.
- **Government and regulatory** — permits, tariffs, public filings, compliance manuals.
- **Manufacturing and procurement** — spec sheets, safety data sheets, bills of material, supplier contracts.
- **Legal** — contracts, amendments, filings, case documents at portfolio scale.

Each of these is a market where someone is currently reading documents by hand and wishing they had the data instead.

How we work with you

Pilot. We take one well-defined slice of your domain, map the schema, and deliver a verified dataset you can inspect field by field. You see exactly what “structured” means for your data before committing to scale. Fixed scope, fixed price.

Scale. Once the pilot proves out, we expand across the full domain: more documents, more regions, more entities, refreshed on your schedule.

Ongoing feed. For sources that change, we run continuous, versioned re-extraction so your dataset stays current without anyone re-keying a thing.

You keep the schema and the data. We keep the engine running.

Why now

The cost of reading a document carefully used to be human time, so structured data at scale was a luxury only the biggest players could afford. That floor just dropped. What took a room of analysts a quarter, the right engine now does in a run, with verification a human process never had.

The winners in every document-heavy industry over the next few years will be the ones who turned their filings into data first. We built the machine that does it. Now we point it at your documents.

Call to action

Have a document domain that should be a database?

Tell us about the documents you wish you could query and the fields you wish you had. We will come back with a proposed schema and a fixed-price pilot scoped to your material.

[[Start a pilot](#)] [[Book a call](#)]

FAQ

How is this different from OCR or a document AI tool? Those give you text or a one-off answer. We give you a defined, verified dataset at the grain you query, with a report of exactly how complete and trustworthy every field is.

How do we know the data is right? Every extraction runs through a verification pass. You receive a gap report showing what was confirmed, what the source never contained, and what needs review. You never have to take an unchecked number on faith.

What if our documents are inconsistent or badly formatted? That is the normal case, not the exception. Inconsistent, blocked, oversized, and irregular sources are what the engine was built to handle.

Our documents are scans, not digital text. Can you still do it? Yes. Scanned, image-only, and photographed documents are handled the same as born-digital PDFs. The output looks identical either way. You cannot tell from the data which source was clean and which was a fax.

How big can the files be? From a one-page invoice to filings that run to hundreds of megabytes. Size is a matter of runtime, not capability, and it never changes the shape of what you get back.

How do you price it? One scoped price per engagement. We look at your material up front and quote a fixed number for the pilot. Whether the documents are text or scanned, tiny or enormous, clean or chaotic, the difficulty is built into the work rather

than passed to you as surprise line items.

Do we own the result? Yes. You own the schema and the data. We provide the extraction engine and the delivery.

Can you handle our volume? The engine fans out across regions and entities and re-runs non-destructively, so scale is a matter of pointing it at more sources, not rebuilding anything.

How do we start? A fixed-scope pilot on one slice of your domain. Low commitment, real output from your own documents.

Structura — the structured layer your documents never had.